

Analyze_ab_test_Szymon_Debbski

March 23, 2021

0.1 Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

0.2 Table of Contents

- Introduction
- Part I - Probability
- Part II - A/B Test
- Part III - Regression

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the [RUBRIC](#).

Part I - Probability

To get started, let's import our libraries.

```
[55]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we
↪ set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset and take a look at the top few rows here:

```
[56]: df = pd.read_csv('ab_data.csv')
```

```
[57]: df.head(5)
```

```
[57]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

b. Use the below cell to find the number of rows in the dataset.

```
[58]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         294478 non-null  int64
1   timestamp       294478 non-null  object
2   group           294478 non-null  object
3   landing_page    294478 non-null  object
4   converted       294478 non-null  int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

c. The number of unique users in the dataset.

```
[59]: len(df.user_id.unique())
```

```
[59]: 290584
```

d. The proportion of users converted.

```
[60]: round(df.converted.mean() * 100, 2)
```

```
[60]: 11.97
```

e. The number of times the `new_page` and `treatment` don't line up.

```
[61]: (df.query('group == "treatment" and landing_page != "new_page")['user_id'].
      ↪count()
+ df.query('group != "treatment" and landing_page == "new_page")['user_id'].
      ↪count())
```

[61]: 3893

f. Do any of the rows have missing values?

```
[62]: df.isnull().sum()
```

```
[62]: user_id          0
      timestamp      0
      group          0
      landing_page    0
      converted      0
      dtype: int64
```

2. For the rows where **treatment** is not aligned with **new_page** or **control** is not aligned with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
[63]: df2 = df.drop(df[(df['group'] == "treatment") & (df['landing_page'] != "new_page")].index)
      df2 = df2.drop(df2[(df2['group'] != "treatment") & (df2['landing_page'] == "new_page")].index)
```

```
[64]: # Double Check all of the correct rows were removed - this should be 0
      df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].shape[0]
```

[64]: 0

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_ids** are in **df2**?

```
[65]: df2.shape[0]
```

[65]: 290585

```
[66]: len(df2.user_id.unique())
```

[66]: 290584

b. There is one **user_id** repeated in **df2**. What is it?

```
[67]: df2[df2.duplicated(['user_id'], keep=False)]
```

```
[67]:   user_id  timestamp      group landing_page  converted
1899   773192  2017-01-09 05:37:58.781806  treatment    new_page         0
2893   773192  2017-01-14 02:55:59.590927  treatment    new_page         0
```

c. What is the row information for the repeat **user_id**?

Information is the same except the timestamp therefore I will remove the second row.

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```
[68]: df2 = df2.drop_duplicates(subset='user_id')
```

4. Use **df2** in the below cells to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
[69]: round(df2.converted.mean() * 100, 2)
```

```
[69]: 11.96
```

b. Given that an individual was in the **control** group, what is the probability they converted?

```
[70]: round(df2.query('group == "control"')['converted'].mean() * 100, 2)
```

```
[70]: 12.04
```

c. Given that an individual was in the **treatment** group, what is the probability they converted?

```
[71]: round(df2.query('group == "treatment"')['converted'].mean() * 100, 2)
```

```
[71]: 11.88
```

d. What is the probability that an individual received the new page?

```
[72]: round(df2.query('landing_page == "new_page"').count()[0]/df2.shape[0] * 100, 2)
```

```
[72]: 50.01
```

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

0.3 Answer:

* The conversion rate is slightly higher for the control group - 12.04 % vs 11.88% treatment group

* However based on this information we cannot conclude with any certainty which page leads to more conversions ### Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a

Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new} , which are the converted rates for the old and new pages.

* **Null hypothesis** - if the p value $\geq 5\%$ the old page has better conversion rate

* **Alternative hypothesis** - if the p value $< 5\%$ the new page has a better conversion rate 2. Assume under the null hypothesis, p_{new} and p_{old} both have “true” success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn’t make complete sense right now, don’t worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

```
[73]: df2.head(1)
```

```
[73]:   user_id      timestamp  group landing_page  converted
0   851104  2017-01-21 22:11:48.556739  control    old_page         0
```

a. What is the **convert rate** for p_{new} under the null?

```
[74]: p_new = df2.converted.mean()
p_new
```

```
[74]: 0.11959708724499628
```

b. What is the **convert rate** for p_{old} under the null?

```
[75]: p_old = df2.converted.mean()
p_old
```

```
[75]: 0.11959708724499628
```

c. What is n_{new} ?

```
[76]: n_new = df2.query('landing_page == "new_page"').count()[0]
n_new
```

```
[76]: 145310
```

d. What is n_{old} ?

```
[77]: n_old = df2.query('landing_page == "old_page"').count()[0]
n_old
```

[77]: 145274

- e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
[78]: new_page_converted = np.random.binomial(1, p_new, n_new)
      new_page_converted.mean()
```

[78]: 0.11807859059940816

- f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
[79]: old_page_converted = np.random.binomial(1, p_old, n_old)
      old_page_converted.mean()
```

[79]: 0.1197599019783306

- g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

```
[80]: #new_page and old_page have different sizes thats why we use means
      new_page_converted.mean() - old_page_converted.mean()
```

[80]: -0.0016813113789224399

- h. Simulate 10,000 $p_{new} - p_{old}$ values using this same process similarly to the one you calculated in parts **a. through g.** above. Store all 10,000 values in a numpy array called **p_diffs**.

```
[81]: new_page_conv = np.random.binomial(n_new, p_new, 10000)/n_new
      old_page_conv = np.random.binomial(n_old, p_old, 10000)/n_old
      p_diffs = new_page_conv - old_page_conv
```

- i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
[82]: plt.hist(p_diffs);
      plt.title('Simulated difference of new_page and old_page conversion rate');
```



j. What proportion of the `p_diffs` are greater than the actual difference observed in `ab_data.csv`?

```
[83]: act_new_conv = df2.query('landing_page == "new_page"')['converted'].mean()
      act_old_conv = df2.query('landing_page == "old_page"')['converted'].mean()

      act_diff = act_new_conv - act_old_conv
      act_diff
```


```
[83]: -0.0015782389853555567
```

```
[84]: p_diffs = np.array(p_diffs)
```

```
[85]: round((p_diffs > act_diff).mean() * 100, 2)
```

```
[85]: 90.62
```

```
[86]: plt.hist(p_diffs);
      plt.axvline(act_diff, c='red');
```



- k. In words, explain what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

In part J we computed the P-Value which in our example is 90%. The greater the value is the more random our results are and therefore NOT significant.

Based on this result we fail to reject the null hypothesis. This means the company should stay with the old page as there is no statistical evidence that the new page has a better conversion rate.

- l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the the number of rows associated with the old page and new pages, respectively.

```
[87]: import statsmodels.api as sm
```

```
convert_old = df2.query('landing_page == "old_page"')['converted'].sum()
convert_new = df2.query('landing_page == "new_page"')['converted'].sum()
n_old = df2.query('landing_page == "old_page"').count()[0]
n_new = df2.query('landing_page == "new_page"').count()[0]

convert_old, n_old, convert_new, n_new
```



```
[87]: (17489, 145274, 17264, 145310)
```

- m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. [Here](#) is a helpful link on using the built in.

```
[88]: z_score, p_value = sm.stats.proportions_ztest([convert_new, convert_old],  
→ [n_new, n_old], alternative='larger')  
z_score, p_value
```

```
[88]: (-1.3109241984234394, 0.9050583127590245)
```

```
[89]: from scipy.stats import norm  
print(norm.ppf(1-(0.05))) #critical value for onesided test at 95%
```

```
1.6448536269514722
```

- n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts **j.** and **k.**?

The p-value and the z_score further prove our previous conclusion that we cannot reject the null hypothesis.

The p-value corresponds to the p-value we computed in part J - which confirms that our results are not significant.

The z-score is less than the critical value (1.64) which also proves that we cannot reject our null hypothesis.

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the previous A/B test can also be achieved by performing regression.

- a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Logistic Regression.

- b. The goal is to use **statsmodels** to fit the regression model you specified in part **a.** to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
[90]: df2.head(1)
```

```
[90]:   user_id      timestamp  group landing_page  converted  
0   851104  2017-01-21 22:11:48.556739  control    old_page         0
```

```
[91]: df2['intercept'] = 1
```

```
df2['ab_page'] = pd.get_dummies(df2['group'])['treatment']

df2.head(1)
```

```
[91]:   user_id          timestamp    group landing_page  converted \
0   851104  2017-01-21 22:11:48.556739  control    old_page        0

      intercept  ab_page
0             1        0
```

- c. Use **statsmodels** to import your regression model. Instantiate the model, and fit the model using the two columns you created in part **b.** to predict whether or not an individual converts.

```
[92]: import statsmodels.api as sm
log_mod = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
```

- d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
[93]: results = log_mod.fit()
results.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.366118
      Iterations 6
```

```
[93]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                Logit Regression Results
=====
Dep. Variable:                  converted    No. Observations:                  290584
Model:                            Logit      Df Residuals:                  290582
Method:                            MLE        Df Model:                        1
Date:                Tue, 23 Mar 2021    Pseudo R-squ.:                  8.077e-06
Time:                12:08:42      Log-Likelihood:                 -1.0639e+05
converged:                            True      LL-Null:                   -1.0639e+05
Covariance Type:            nonrobust      LLR p-value:                   0.1899
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept    -1.9888      0.008   -246.669      0.000      -2.005      -1.973
ab_page      -0.0150      0.011    -1.311      0.190      -0.037      0.007
=====
      """
```

- e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint:** What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in the **Part II**?

The p-value for the above model is 0.19 which is greater than 5% which indicates that the alternative hypothesis is not significant.

When comparing this with the model in Part II the conclusion is the same - that the conversion rate for the old_page is the same or better.

One thing to note that the p-value is different than in part II.

- f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

It is a good idea to consider other factors in the model. However, there is always the risk the model will become too complex and it will be hard to come up with any clear conclusions.

- g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. [Here](#) are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables.** Provide the statistical output as well as a written response to answer this question.

```
[94]: countries_df = pd.read_csv('./countries.csv')
df_new = countries_df.set_index('user_id').join(df2.set_index('user_id'),
        ↪how='inner')
df_new.head(1)
```

```
[94]:      country      timestamp  group landing_page  converted \
user_id
834778      UK  2017-01-14 23:08:43.304998  control    old_page         0

      intercept  ab_page
user_id
834778         1         0
```

```
[95]: df_new.groupby(['country'])['converted'].mean()
```

```
[95]: country
CA    0.115318
UK    0.120594
US    0.119547
Name: converted, dtype: float64
```

```
[96]: df_new.country.unique()
```

```
[96]: array(['UK', 'US', 'CA'], dtype=object)
```

```
[97]: df_new[['UK', 'US', 'CA']] = pd.get_dummies(df_new['country'])
df_new.head(5)
```

```
[97]:
```

	country	timestamp	group	landing_page \
user_id				
834778	UK	2017-01-14 23:08:43.304998	control	old_page
928468	US	2017-01-23 14:44:16.387854	treatment	new_page
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page
711597	UK	2017-01-22 03:14:24.763511	control	old_page
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

	converted	intercept	ab_page	UK	US	CA
user_id						
834778	0	1	0	0	1	0
928468	0	1	1	0	0	1
822059	1	1	1	0	1	0
711597	0	1	0	0	1	0
710616	0	1	1	0	1	0

```
[98]: df_new['intercept'] = 1

mlr = sm.Logit(df_new['converted'], df_new[['intercept', 'ab_page', 'UK', 'CA']])
results = mlr.fit()
results.summary()
```

Optimization terminated successfully.
Current function value: 0.366113
Iterations 6

```
[98]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                Logit Regression Results
=====
Dep. Variable:                converted    No. Observations:                290584
Model:                        Logit        Df Residuals:                290580
Method:                        MLE          Df Model:                    3
Date:                        Tue, 23 Mar 2021    Pseudo R-squ.:                2.323e-05
Time:                        12:08:49      Log-Likelihood:                -1.0639e+05
converged:                    True          LL-Null:                    -1.0639e+05
Covariance Type:              nonrobust      LLR p-value:                0.1760
=====

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-1.9794	0.013	-155.415	0.000	-2.004	-1.954
ab_page	-0.0149	0.011	-1.307	0.191	-0.037	0.007
UK	-0.0506	0.028	-1.784	0.074	-0.106	0.005
CA	-0.0099	0.013	-0.743	0.457	-0.036	0.016

```
=====
"""
```

Looking at the results we can conclude that the country does not impact the conversion rate as the p-value for both cases is $> 5\%$.

- h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

```
[99]: df_new['UK_ab_page'] = df_new['UK'] * df_new['ab_page']
df_new.head(1)
```

```
[99]:      country      timestamp  group landing_page  converted  \
user_id
834778      UK  2017-01-14 23:08:43.304998  control    old_page        0

      intercept  ab_page  UK  US  CA  UK_ab_page
user_id
834778         1        0  0  1  0          0
```

```
[100]: df_new['CA_ab_page'] = df_new['CA'] * df_new['ab_page']
df_new.head(1)
```

```
[100]:      country      timestamp  group landing_page  converted  \
user_id
834778      UK  2017-01-14 23:08:43.304998  control    old_page        0

      intercept  ab_page  UK  US  CA  UK_ab_page  CA_ab_page
user_id
834778         1        0  0  1  0          0          0
```

```
[101]: df_new['intercept'] = 1

mlr = sm.Logit(df_new['converted'], df_new[['intercept', 'ab_page', 'UK', 'CA',
      ↪ 'UK_ab_page', 'CA_ab_page']])
results = mlr.fit()
results.summary()
```

Optimization terminated successfully.

Current function value: 0.366109
Iterations 6

```
[101]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
```

```

Dep. Variable:          converted    No. Observations:          290584
Model:                  Logit        Df Residuals:              290578
Method:                  MLE         Df Model:                  5
Date:                   Tue, 23 Mar 2021    Pseudo R-squ.:            3.482e-05
Time:                   12:08:52          Log-Likelihood:           -1.0639e+05
converged:              True          LL-Null:                  -1.0639e+05
Covariance Type:        nonrobust       LLR p-value:              0.1920

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-1.9922	0.016	-123.457	0.000	-2.024	-1.961
ab_page	0.0108	0.023	0.475	0.635	-0.034	0.056
UK	-0.0118	0.040	-0.296	0.767	-0.090	0.066
CA	0.0057	0.019	0.306	0.760	-0.031	0.043
UK_ab_page	-0.0783	0.057	-1.378	0.168	-0.190	0.033
CA_ab_page	-0.0314	0.027	-1.181	0.238	-0.084	0.021

```

=====
"""

```

Conclusions

Based on the results above and the interactions between country and the ab_page we can conclude that none of the variables have significant p-values. Having this in mind we would fail to reject the null hypothesis.

Taking into account all the results from all the tests my conclusion is that we should stick with the old page as there is no evidence that the new page is better. Moreover the old page seems to have a slightly better conversion rate.